



Automatic LLM Benchmark Analysis for Text2SQL



Authors: Cristian DEGNI, Francesco GIANNUZZO

Supervised by: Prof. Paolo PAPOTTI, Simone PAPICCHIO (PhD)



Project Background & Motivation

🧠 **LLMs** and approaches like **Text-to-SQL** and **Text-to-Code** are transforming how we query databases, especially in complex, enterprise settings.

📊 Existing **public benchmarks** often **fail** to **reflect** the **diversity and complexity** of **real-world datasets**, limiting accurate model evaluation.

It is therefore important to:

🔍 **Evaluate** the **robustness** and **generalization** of models on SQL-centric tasks.

⚖️ Compare **Text-to-Code** systematically **against** established methods like **Text-to-SQL**.

Project Objectives

- 1) **Develop a web application using GRADIO to make QATCH more user-friendly and accessible for end users.**
- 2) **Compare the Text-to-SQL and Text-to-Code approaches using different LMMs on different databases using Natural Language Question for querying the data.**
- 3) **Evaluate models performance using QATCH metrics for multiple scenarios and Code metrics for the Text-To-Code task.**
- 4) **Analyze the results obtained from models on different SQL-centric and non-SQL datasets to compare the effectiveness of both task approaches under realistic conditions.**

Workflow Overview

It generates **SQL queries**, **Python scripts** or **answer**



Predictor









Executor



Evaluator



PREDICTOR



EVALUATOR

EXECUTOR

Workflow Overview



It runs those **queries** or **scripts** to produce **result tables**.



EXECUTOR

Load Data Files

Data

Creation Virtual
SandBox

E2B

API E2B

Executor

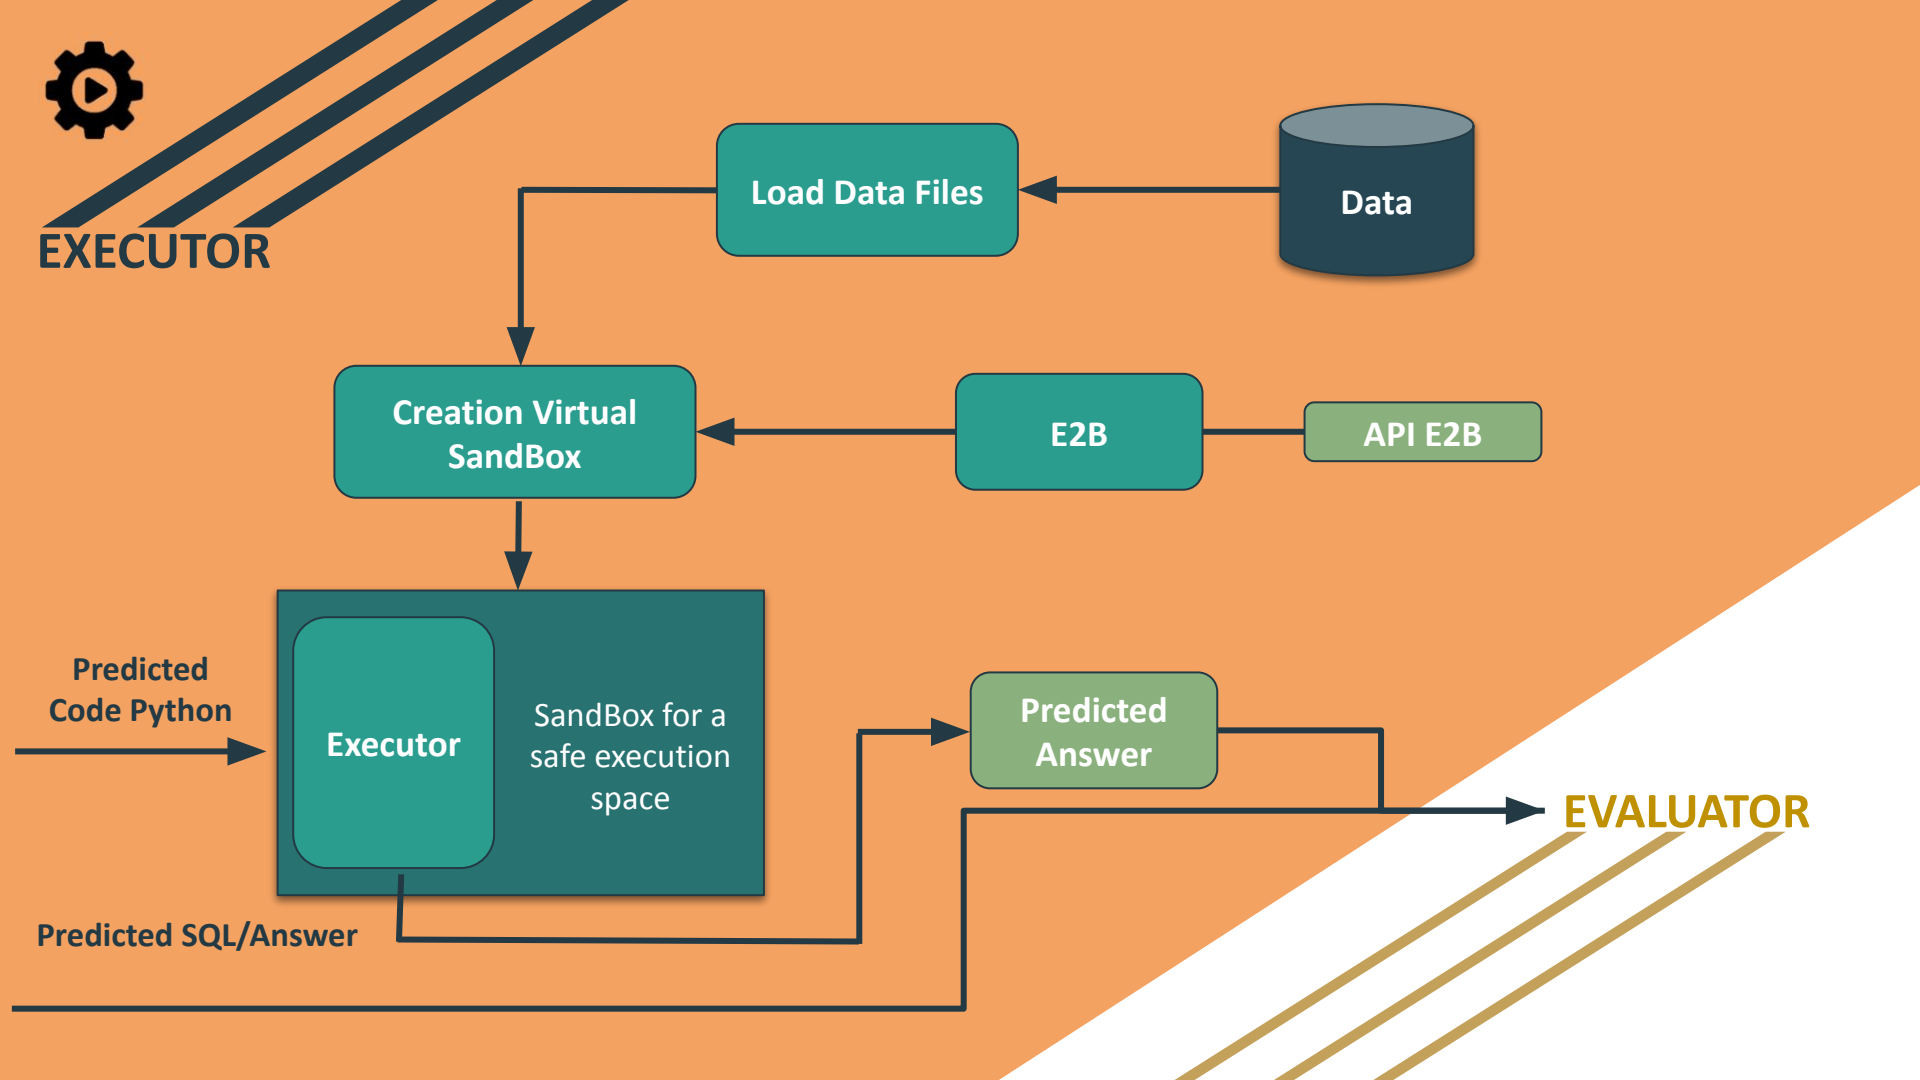
SandBox for a
safe execution
space

Predicted
Answer


Predicted
Code Python

Predicted SQL/Answer

EVALUATOR



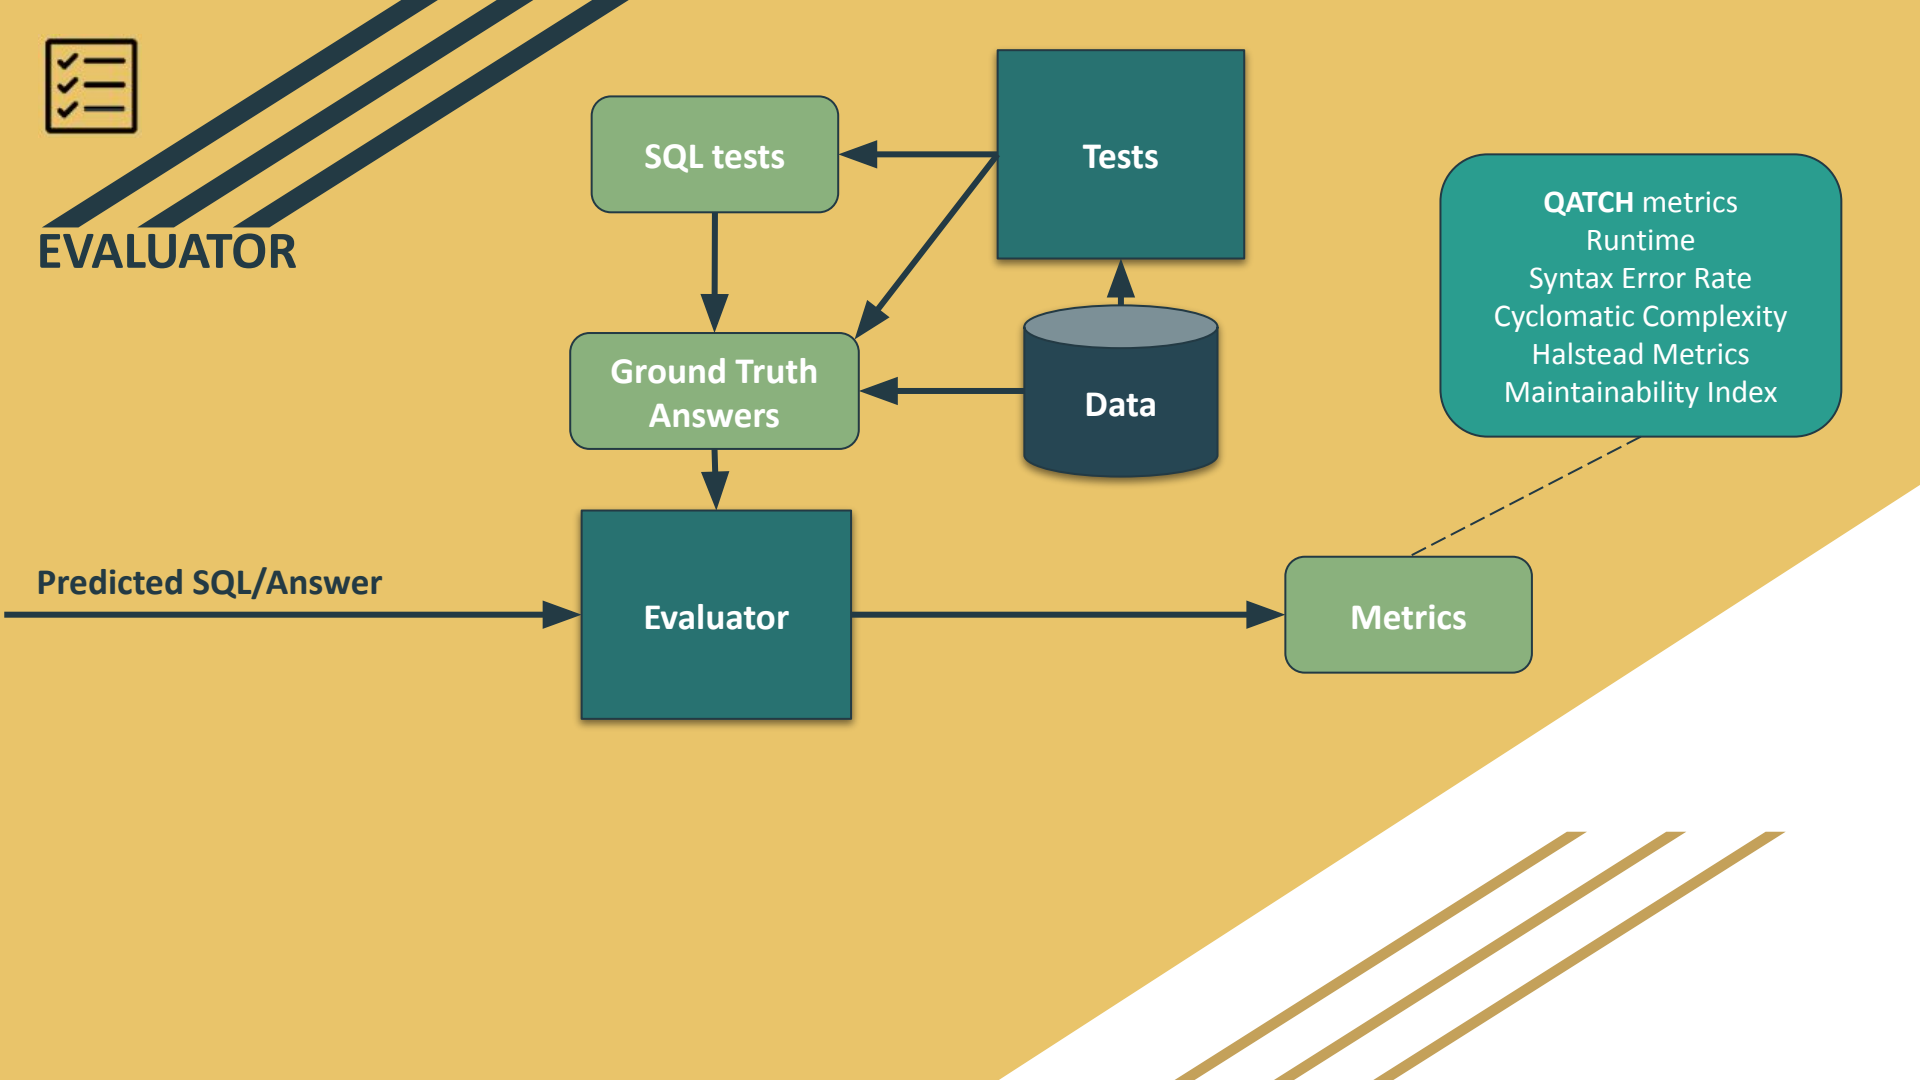
Workflow Overview



It compares the **ground truth** tables with predictions for computes **performance metrics**.



EVALUATOR



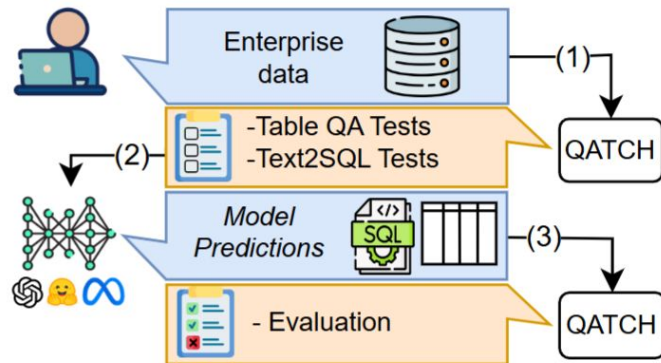
SQL-centric (Text2SQL & Table QA)

Text-to-SQL:

Goal : translate an **NL question** and **data context** into a **SQL query**, which is **executed** to **obtain table results**. **QATCH evaluates** execution accuracy and fine-grained metrics for a full performance profile.

Table Question Answering:

LLM takes an **NL question** and **data context** to produce **directly** an **answer**, without issuing SQL engine. **QATCH** then **compares** this **output** to the **ground-truth answer** to **measure accuracy** and **pinpoint errors** at the cell and tuple levels.



Text2SQL

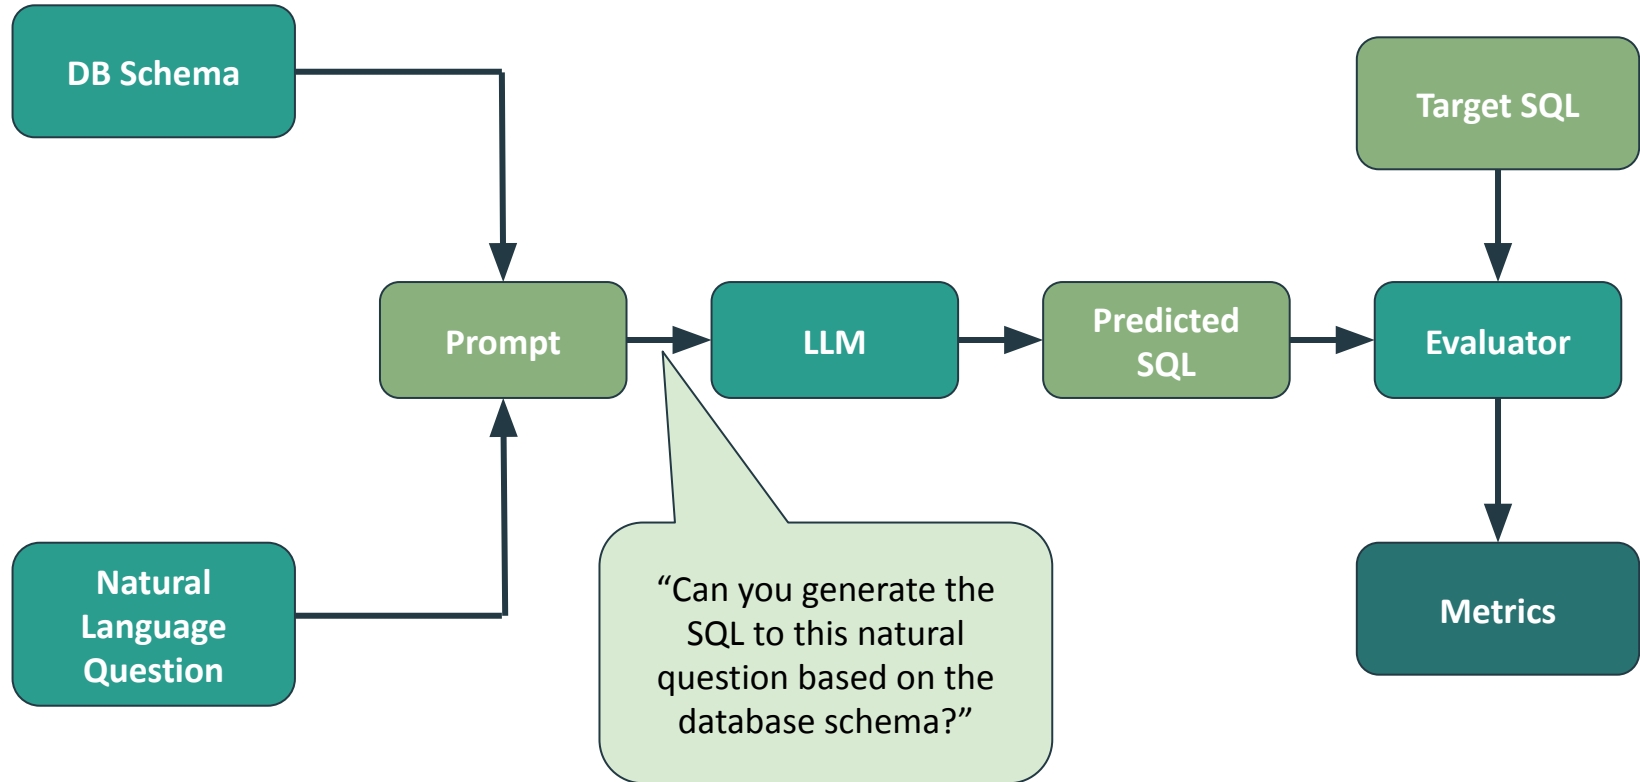
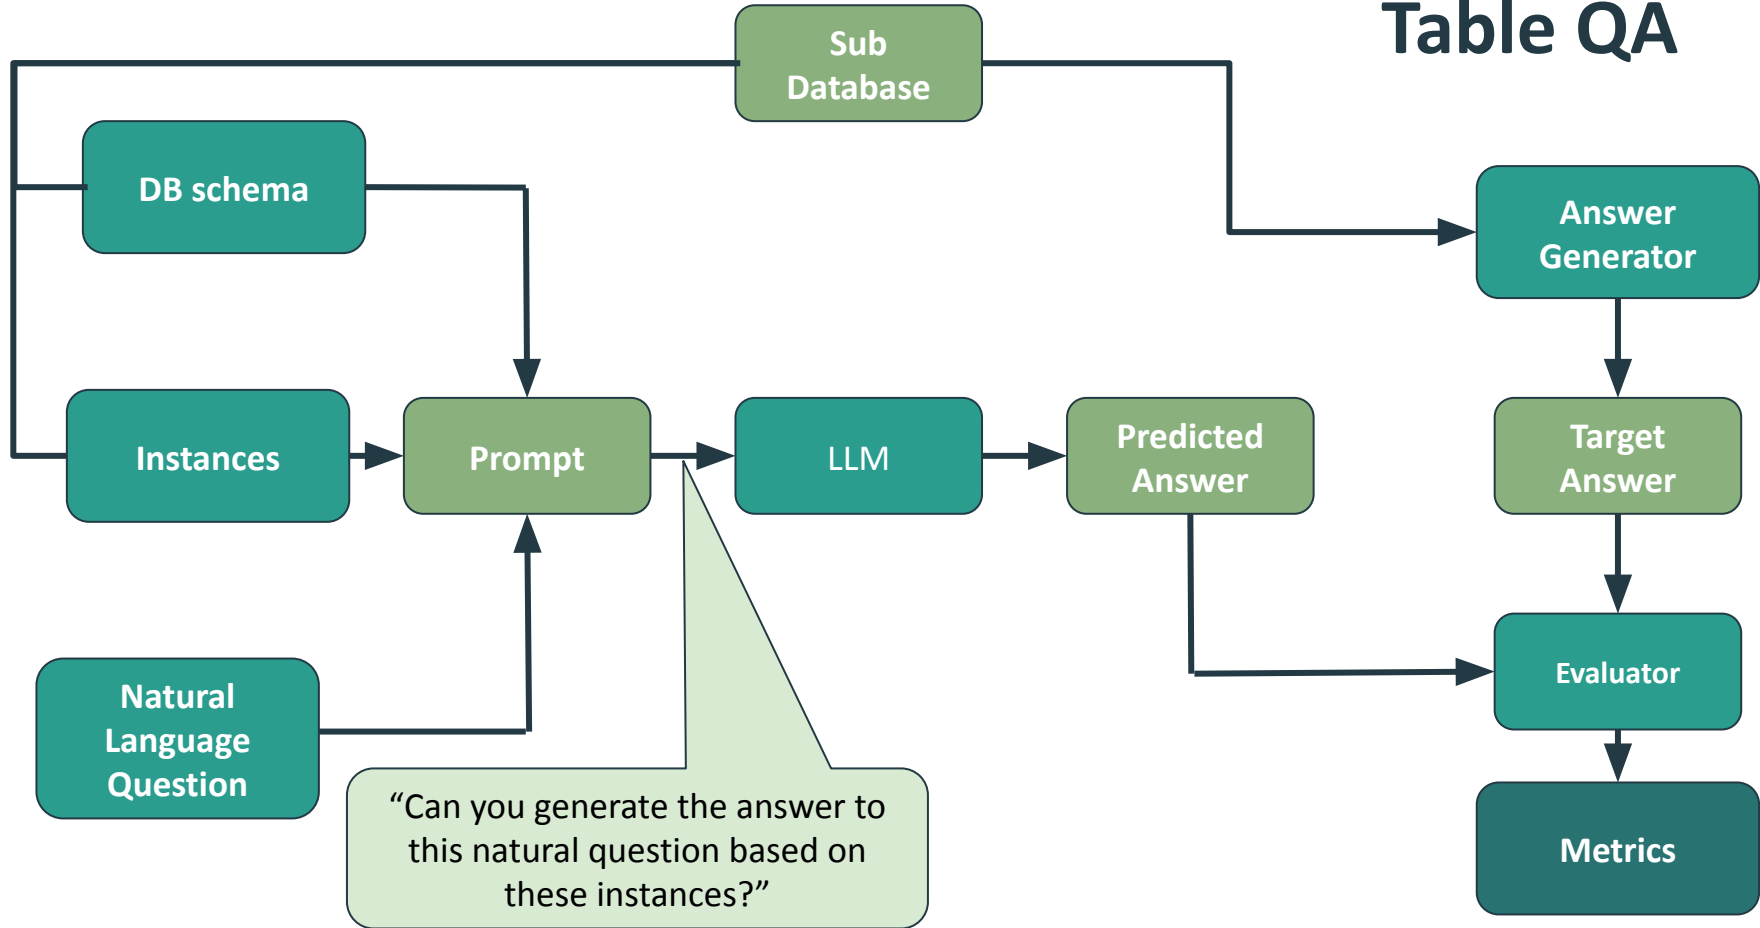


Table QA



Text2Code

Definition:

Map a natural language question and database context into executable Python code that returns the answer table.

Approaches:

Vanilla: single-prompt for directly generate a Python code

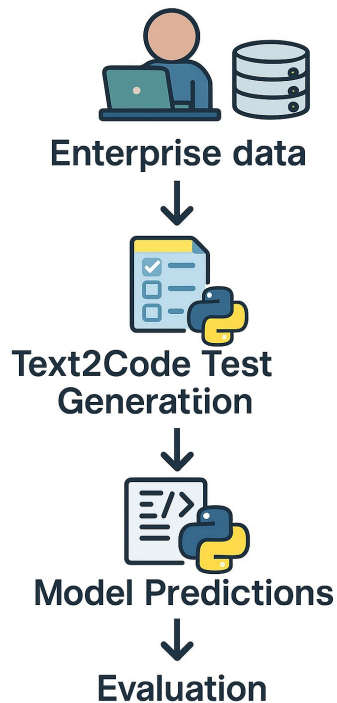
CodexDB: LLM-driven pseudocode translated into Python without leaking gold SQL

Binder: two-stage LLM decomposition with intermediate mapping

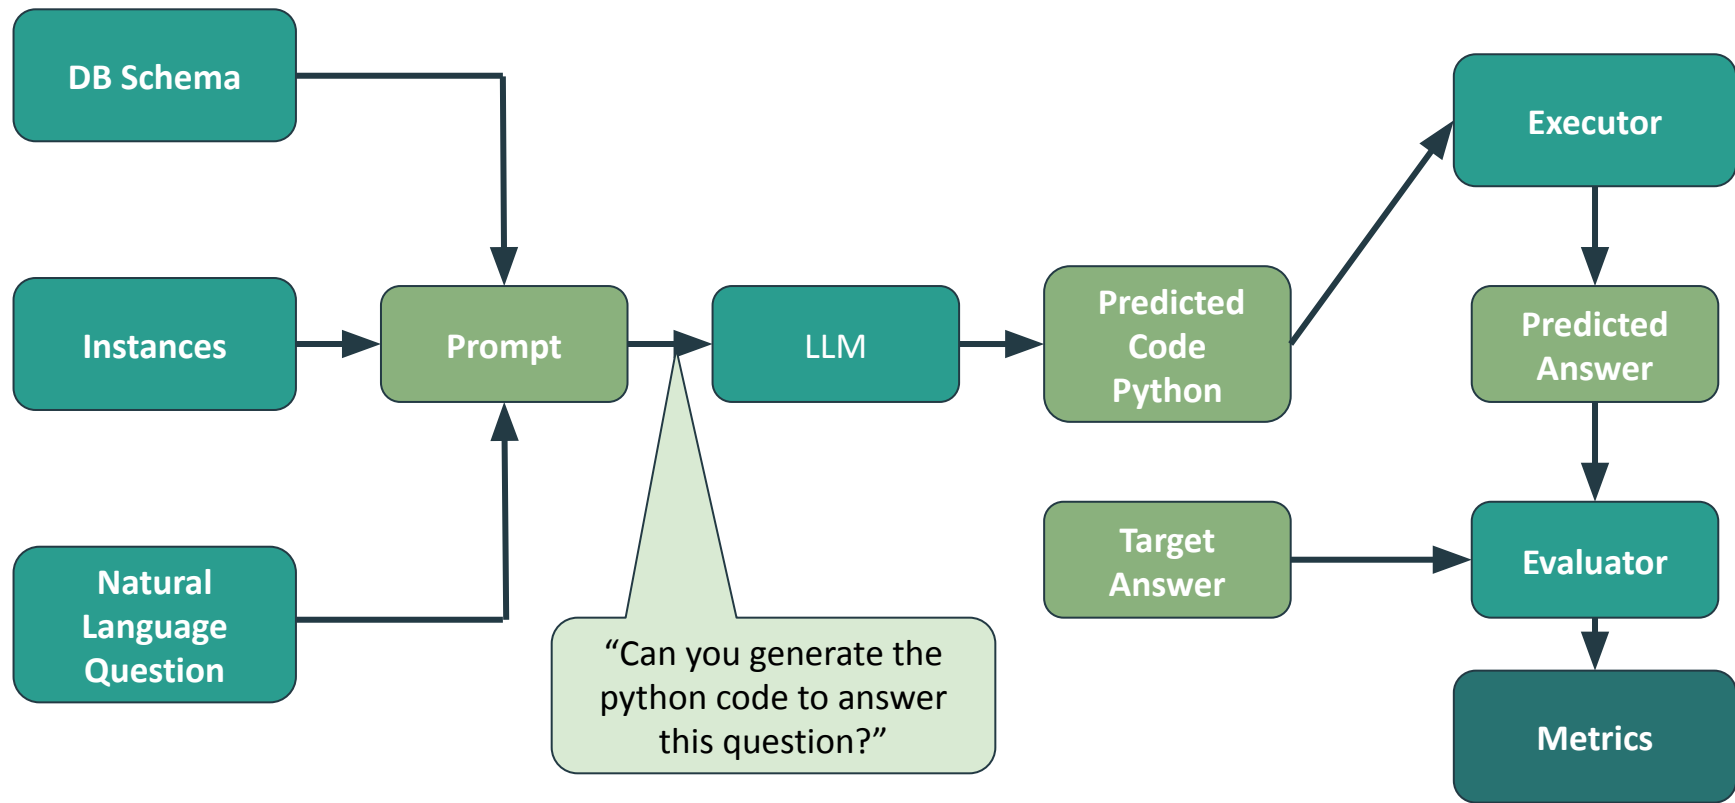


Quality code metrics:

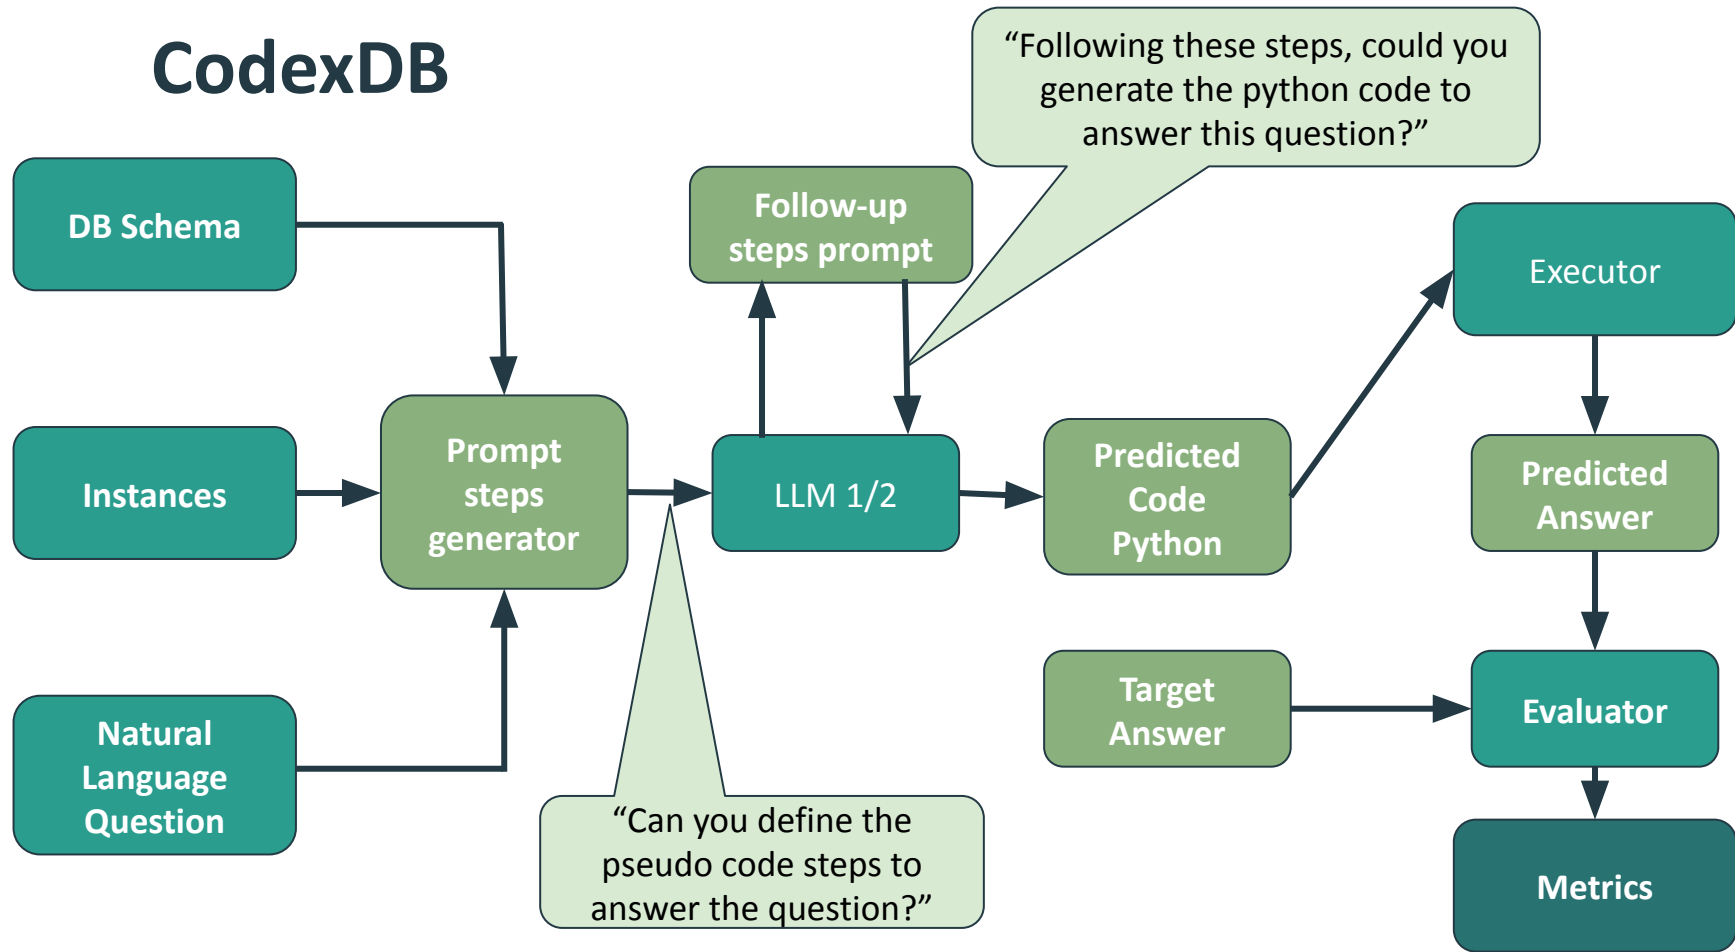
Cyclomatic Complexity, Halstead Measures, Maintainability Index



Vanilla Code Python

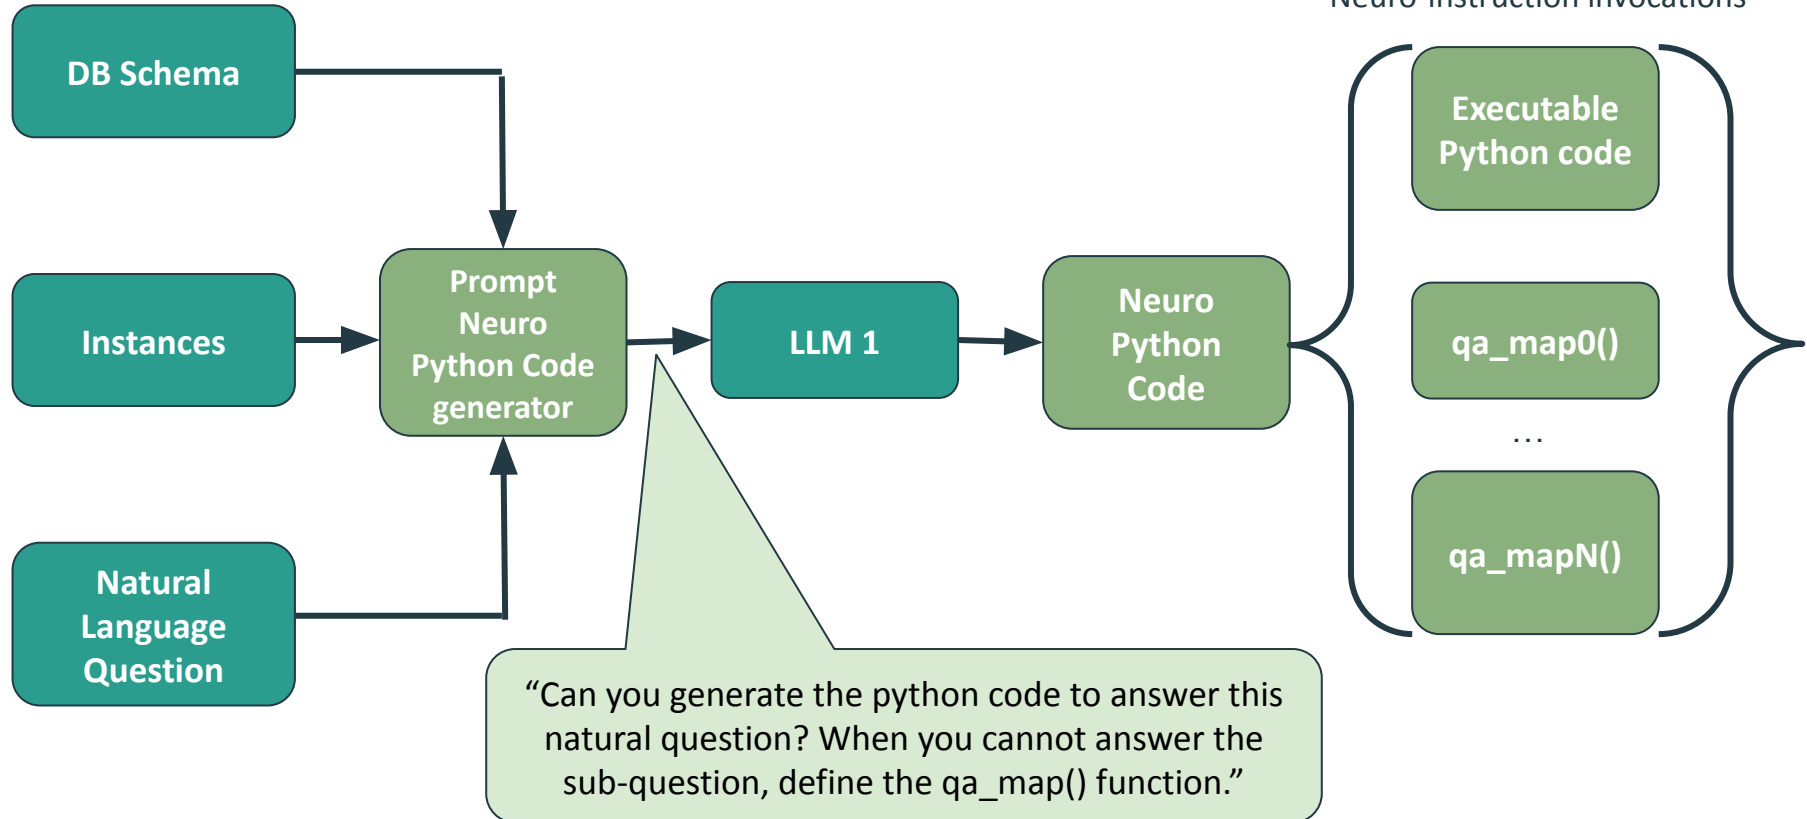


CodexDB

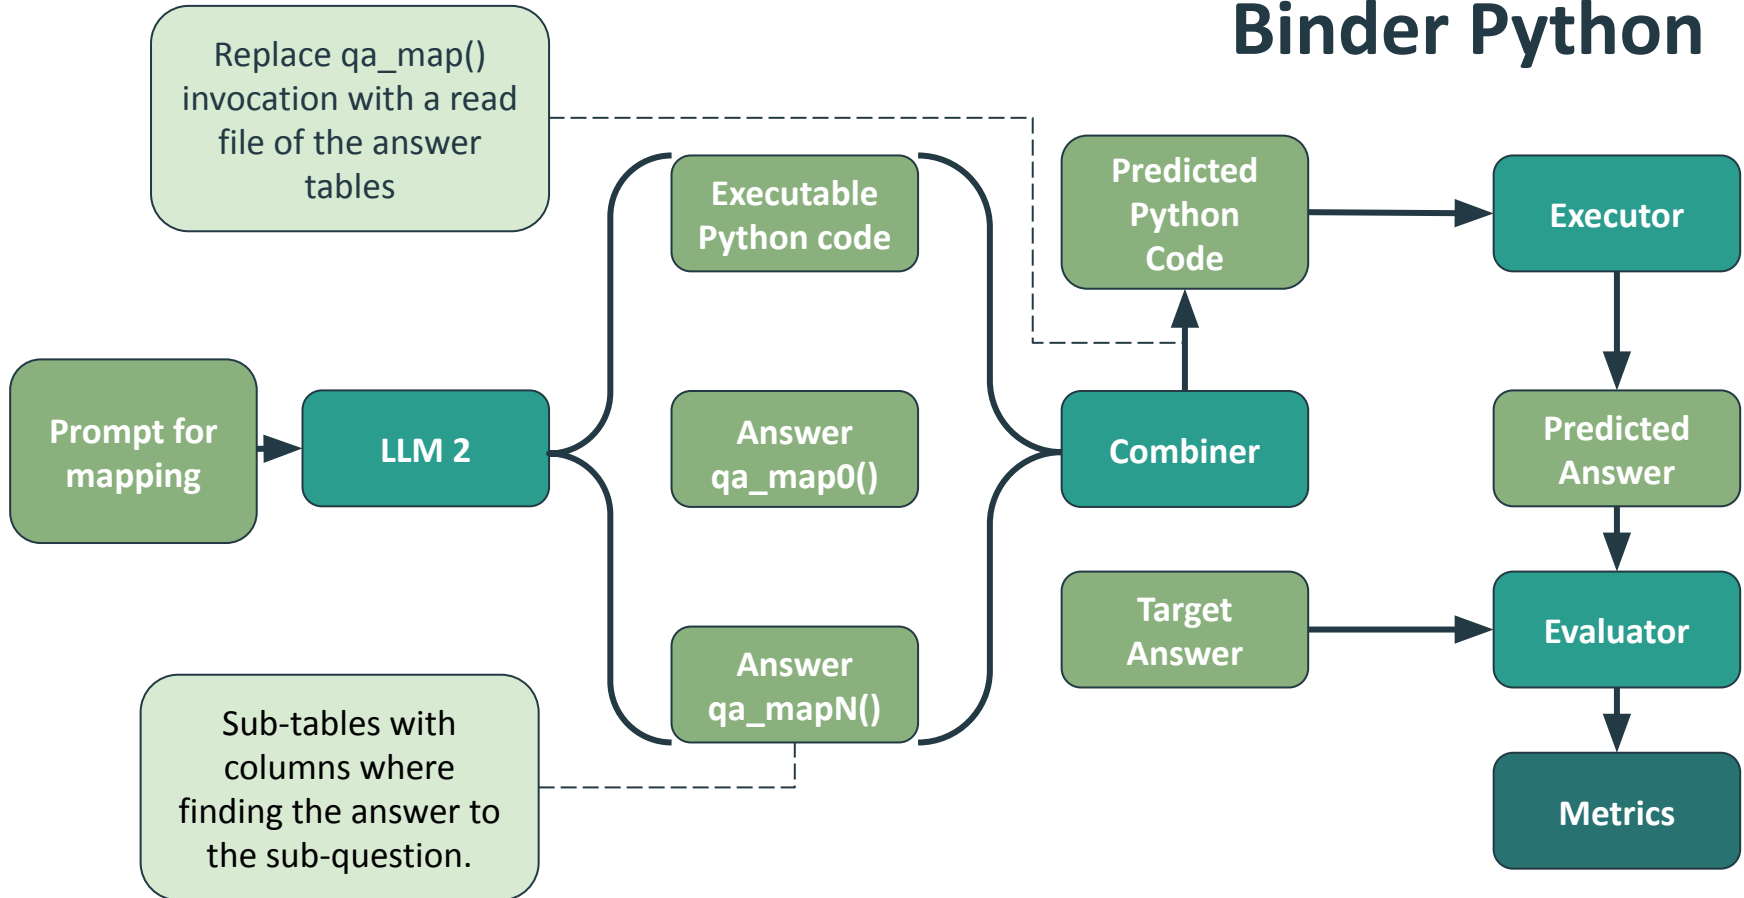


Binder Python

Neuro-instruction invocations

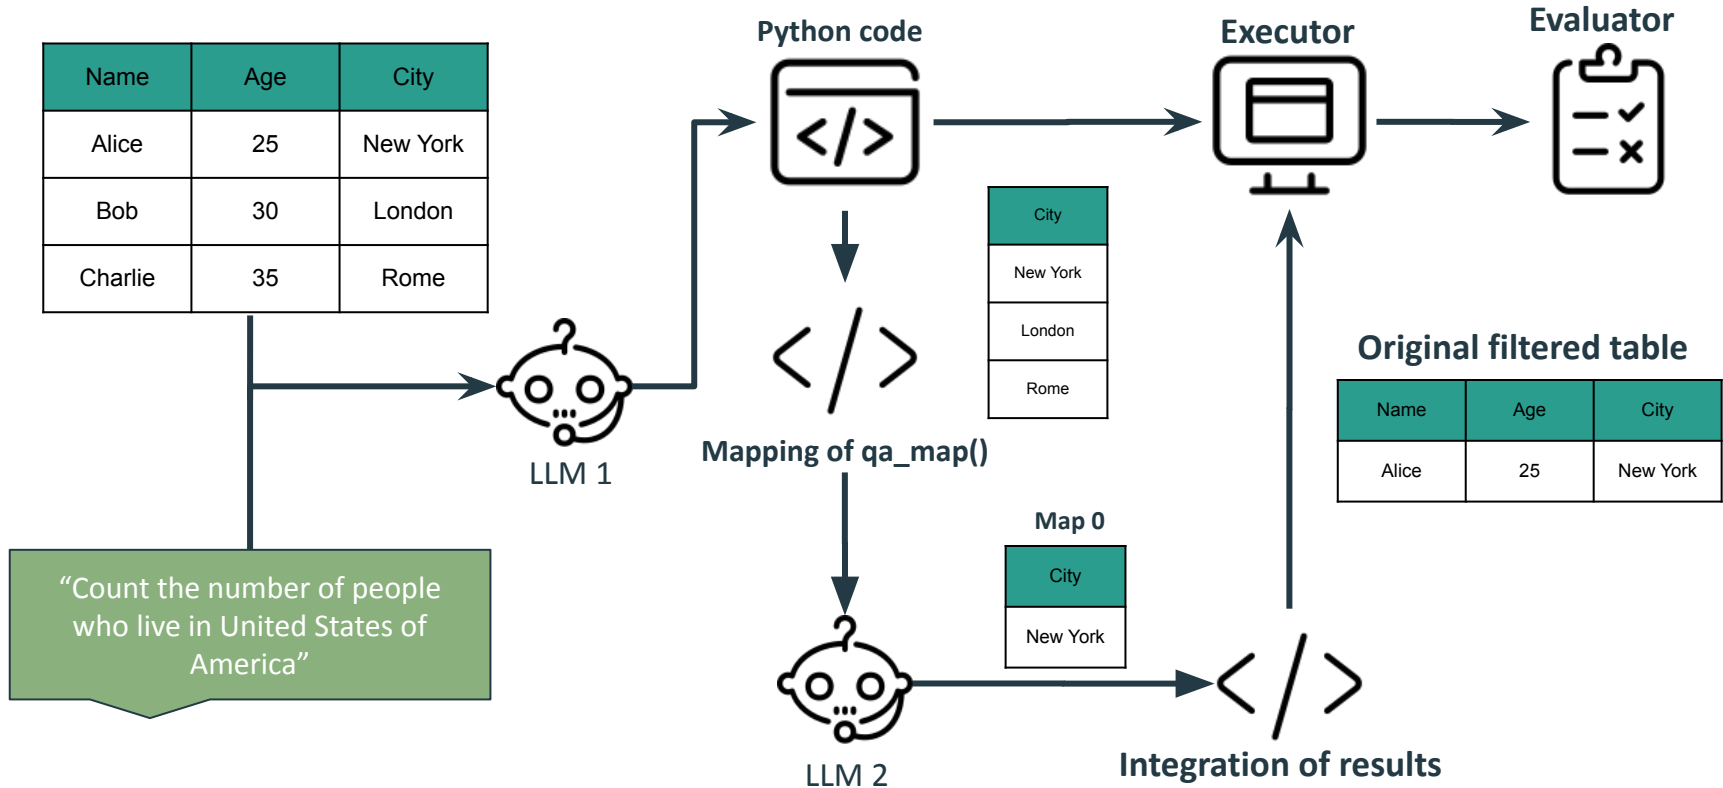


Binder Python



Binder Python example

`qa_map('table_default.csv', 'which cities are in the United States', ['City'])`



Binder Python example

Neuro Python Code generated

```
# Use qa_map to filter the table
for cities in the United States
filtered_table =
qa_map('table_default.csv', 'which
cities are in the United States',
['City'])

# Count the number of people who
live in the United States
count_usa = len(filtered_table)

# Print the result
print([[count_usa]])
```



Code after results integrations

```
#Use qa_map to filter the table
for cities in the United States
filtered_table =
pd.read_csv('map0.csv')

# Count the number of people who
live in the United States
count_usa = len(filtered_table)

# Print the result
print([[count_usa]])
```

Metrics: Evaluating Model Performance

To **provide a comprehensive analysis** of the **performance of our models**, we used **two different classes of metrics**:

QATCH Metrics:

Unified evaluation metrics used across QA, Text-to-SQL, and Text-to-Code, measuring both **semantic correctness** and **structural fidelity** of the model outputs.

Code Metrics:

A custom suite designed for Text-to-Code, evaluating the **complexity**, **maintainability**, and **readability** of generated Python code.

Code Metrics

Cyclomatic Complexity

Measures the number of independent “paths” that can be taken in the code execution flow. It is related to conditional logic (if, for, while, try, etc.).

Components

E = arcs,

N = nodes,

P = connected components

$$CC = E - N + 2P$$

Returns a list of values representing the complexity for each independent “path”

Code Metrics

Halstead Metrics

Measures the cognitive complexity of code based on the operators and operands used. It is useful for estimating how difficult it is to understand or maintain a program.

Components

n1: distinct number of operands

n2: distinct number of operators

N1: total number of operands

N2: total number of operators

$$\text{Volume (V)} = (N1 + N2) \times \log_2(n1 + n2)$$

→ amount of information

$$\text{Difficulty (D)} = (n1 / 2) \times (N2 / n2)$$

→ difficulty of comprehension

$$\text{Effort (E)} = D \times V$$

→ estimated mental effort

Code Metrics

Maintainability Index

A composite index that measures how easy a piece of code is to maintain.

Components






CC = Cyclomatic Complexity

HV = Halstead Volume

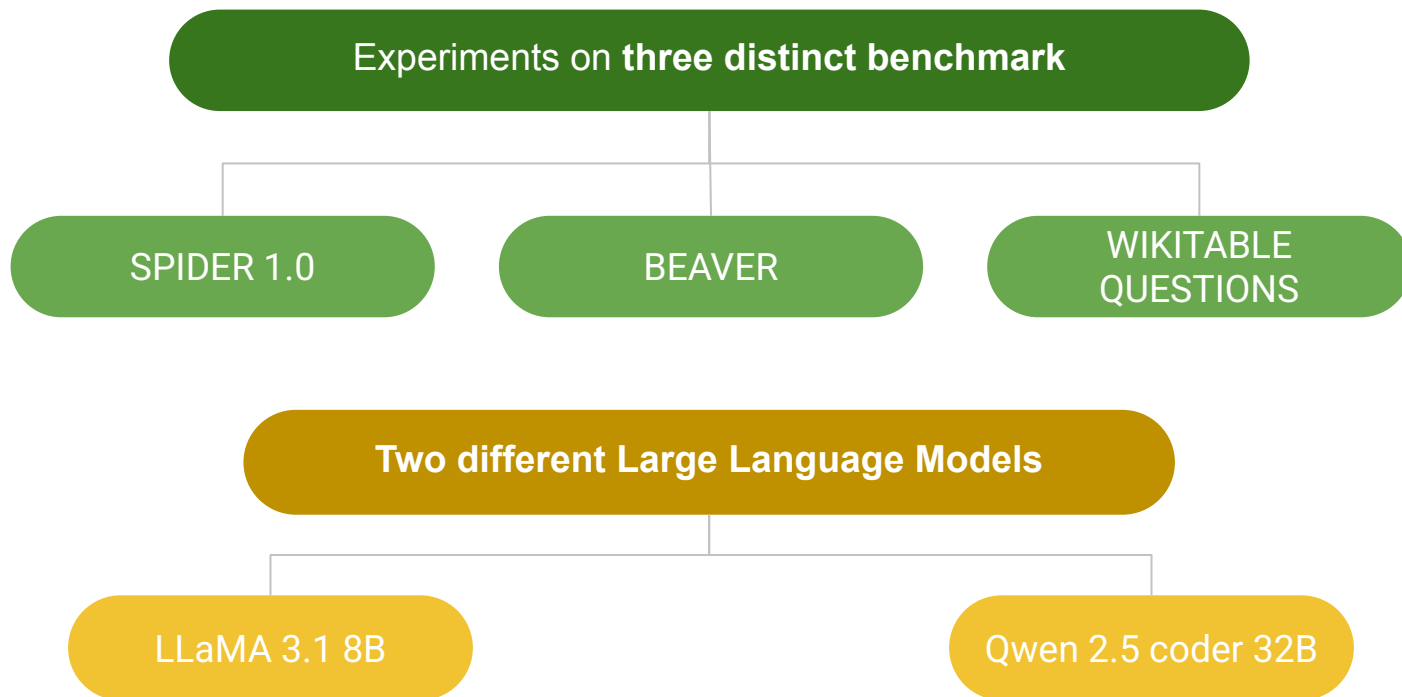
NLC = Number of Lines of Code

$$MI = 171 - 5.2 * \ln(HV) - 0.23 * CC - 16.2 * \ln(NLC)$$

QATCH Metrics


 Cell Precision	Percentage of predicted table cells that are correct.
 Cell Recall	Percentage of ground-truth table cells that were successfully retrieved.
 Tuple Constraint	Exact match on schema, cardinality, and cell values (1 if identical, 0 otherwise).
 Tuple Cardinality	Ratio of predicted to ground-truth tuple counts.
 Tuple Order	Correlation between predicted vs. true tuple ordering (for ORDER BY queries).

Results




Preliminary Tests


[Non-proprietary] Spider avg metrics (Concert Singer)



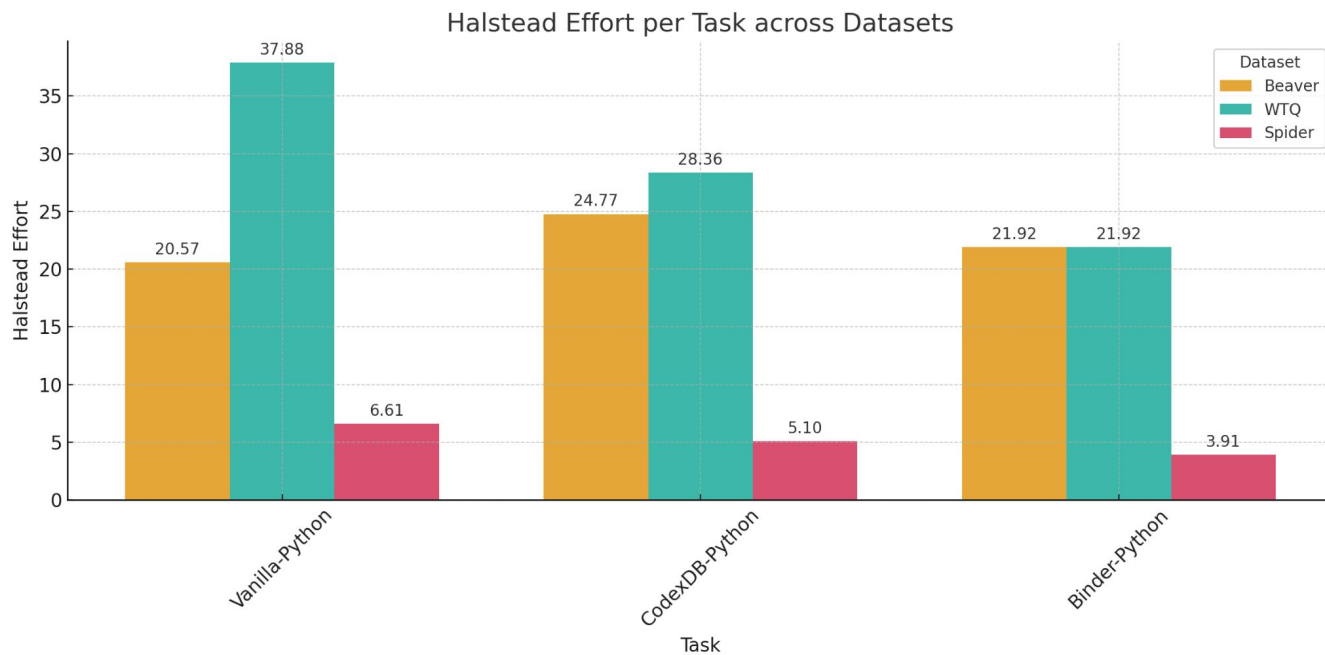
[Proprietary] BEAVER avg metrics (TIME QUARTERS)



QATCH Results



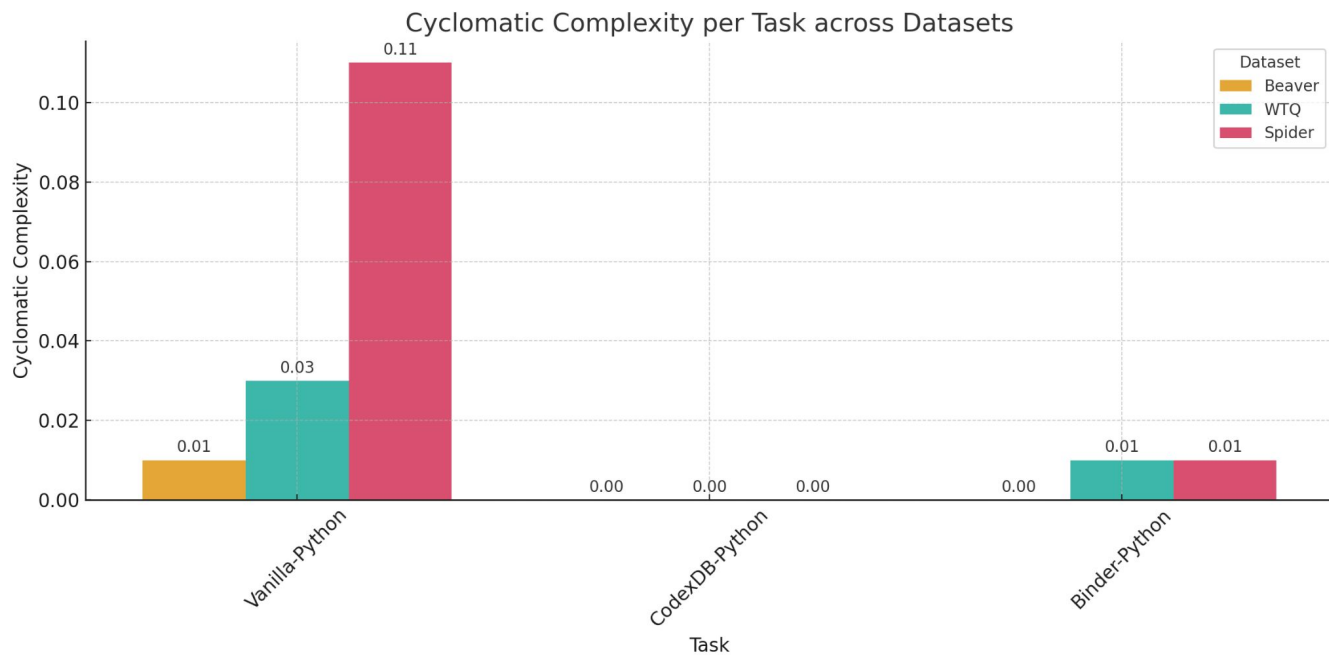
Halstead (effort) Results



Maintainability Index Results



Cyclomatic Complexity Results



Conclusions



Some models show stronger **task-specific reasoning**, especially in structured domains.



LLMs prefer **simpler tasks with clearer structure**, especially when schema info is limited.



Text2SQL is stronger on SQL-centric tasks, while **Text2Code** excels in reasoning intensive ones.



There is **room for improvement** in handling complex datasets and ambiguous queries.



Model performance varies depending on dataset type (e.g., BEAVER vs. SPIDER).



Prompt design and instance diversity play a critical role in code generation quality.

Thanks for your attention ;)

Questions?

Future works



Enhance existing pipelines

Refine Vanilla, CodexDB, and Binder to handle prompt sensitivity and memory issues.



Explore new Text2Code paradigms

Explore methods like chain-of-thought and program-by-example for better reasoning.



Automatic schema induction

Develop automatic schema induction for semi-structured tables like WTQ.



Benchmark on no-SQL-centric datasets

Evaluate performance on datasets like TabFact or HybridQA.



Evaluate different model types

Use code-specialized and multimodal models to assess task suitability.



Scale Table QA to full datasets

Implement full-dataset QA with retrieval and streaming for large tables.